

# CMTH642 - Data Analytics: Advanced Methods

## Assignment 2

Assignment 2 is worth 15% of the final grade. Submit the ipynb file and a generated output file (PDF or HTML). Failing to submit both files will be subject to a mark deduction.

Your output file should include all the tables, plots, and requested figures/values.

Printing the lengthy outputs (e.g., the whole data frame or a list with more than 100 elements) will have a deduction of 5 points. Instead, please use `head()` or `tail()` to have a neat output.

If you preprocess the data in a question, you should continue with the same dataset in the following questions unless otherwise mentioned.

### Preparation:

The dataset is related to Portuguese "Vinho Verde" wines. For more info:

<https://archive.ics.uci.edu/ml/datasets/Wine+Quality>

Import the following files: <http://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/winequality-white.csv>

<http://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/winequality-red.csv>

### Question 1

Join the red wine and white wine datasets by adding the rows of one to the other. Assign the joined data to a data frame and name it `wine_data_all`. (2 points)

The following questions will be answered for `wine_data_all`.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

import seaborn as sns

red_wine_df = pd.read_csv('http://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/winequality-red.csv',
                          sep=';')
white_wine_df = pd.read_csv('http://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/winequality-white.csv',
                             sep=';')

red_wine_df.head()
```

	fixed acidity	volatile acidity	citric acid	residual sugar
chlorides \				
0	7.4	0.70	0.00	1.9

```

0.076
1          7.8          0.88          0.00          2.6
0.098
2          7.8          0.76          0.04          2.3
0.092
3          11.2         0.28          0.56          1.9
0.075
4          7.4          0.70          0.00          1.9
0.076

```

```

      free sulfur dioxide  total sulfur dioxide  density    pH  sulphates
\
0          11.0          34.0    0.9978  3.51      0.56
1          25.0          67.0    0.9968  3.20      0.68
2          15.0          54.0    0.9970  3.26      0.65
3          17.0          60.0    0.9980  3.16      0.58
4          11.0          34.0    0.9978  3.51      0.56

```

```

      alcohol  quality
0         9.4        5
1         9.8        5
2         9.8        5
3         9.8        6
4         9.4        5

```

```
white_wine_df.tail()
```

```

      fixed acidity  volatile acidity  citric acid  residual sugar
chlorides \
4893          6.2          0.21          0.29          1.6
0.039
4894          6.6          0.32          0.36          8.0
0.047
4895          6.5          0.24          0.19          1.2
0.041
4896          5.5          0.29          0.30          1.1
0.022
4897          6.0          0.21          0.38          0.8
0.020

```

```

      free sulfur dioxide  total sulfur dioxide  density    pH
sulphates \
4893          24.0          92.0    0.99114  3.27
0.50
4894          57.0          168.0    0.99490  3.15

```

```

0.46
4895          30.0          111.0  0.99254  2.99
0.46
4896          20.0          110.0  0.98869  3.34
0.38
4897          22.0          98.0  0.98941  3.26
0.32

```

```

      alcohol  quality
4893     11.2        6
4894      9.6        5
4895      9.4        6
4896     12.8        7
4897     11.8        6

```

```

wine_data_all = pd.concat([red_wine_df, white_wine_df],
ignore_index=True)

```

```

wine_data_all.tail()

```

```

      fixed acidity  volatile acidity  citric acid  residual sugar
chlorides \
6492          6.2          0.21          0.29          1.6
0.039
6493          6.6          0.32          0.36          8.0
0.047
6494          6.5          0.24          0.19          1.2
0.041
6495          5.5          0.29          0.30          1.1
0.022
6496          6.0          0.21          0.38          0.8
0.020

```

```

      free sulfur dioxide  total sulfur dioxide  density  pH
sulphates \
6492          24.0          92.0  0.99114  3.27
0.50
6493          57.0          168.0  0.99490  3.15
0.46
6494          30.0          111.0  0.99254  2.99
0.46
6495          20.0          110.0  0.98869  3.34
0.38
6496          22.0          98.0  0.98941  3.26
0.32

```

```

      alcohol  quality
6492     11.2        6
6493      9.6        5
6494      9.4        6

```

6495	12.8	7
6496	11.8	6

## Question 2

Check the data types of the attributes. (2 points)

```
wine_data_all.dtypes
```

```
fixed acidity      float64
volatile acidity   float64
citric acid        float64
residual sugar     float64
chlorides          float64
free sulfur dioxide float64
total sulfur dioxide float64
density            float64
pH                 float64
sulphates          float64
alcohol            float64
quality            int64
dtype: object
```

```
wine_data_all.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 6497 entries, 0 to 6496
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	fixed acidity	6497 non-null	float64
1	volatile acidity	6497 non-null	float64
2	citric acid	6497 non-null	float64
3	residual sugar	6497 non-null	float64
4	chlorides	6497 non-null	float64
5	free sulfur dioxide	6497 non-null	float64
6	total sulfur dioxide	6497 non-null	float64
7	density	6497 non-null	float64
8	pH	6497 non-null	float64
9	sulphates	6497 non-null	float64
10	alcohol	6497 non-null	float64
11	quality	6497 non-null	int64

```
dtypes: float64(11), int64(1)
```

```
memory usage: 609.2 KB
```

## Question 3

Are there any missing values in the dataset? How many? You should not print the whole dataset. (2 points)

```
wine_data_all.isnull().sum()
```

```
fixed acidity      0
volatile acidity   0
citric acid        0
residual sugar     0
chlorides          0
free sulfur dioxide 0
total sulfur dioxide 0
density           0
pH                0
sulphates         0
alcohol           0
quality           0
dtype: int64
```

There are no missing values in the dataset.

## Question 4

What is the correlation between the attributes other than Quality? (8 points)

```
# Select only the columns that are not 'quality'
attributes_other_than_quality = wine_data_all.drop('quality', axis=1)

# Calculate the correlation
correlation_matrix = attributes_other_than_quality.corr()

print(correlation_matrix)
```

	fixed acidity	volatile acidity	citric acid	\
fixed acidity	1.000000	0.219008	0.324436	
volatile acidity	0.219008	1.000000	-0.377981	
citric acid	0.324436	-0.377981	1.000000	
residual sugar	-0.111981	-0.196011	0.142451	
chlorides	0.298195	0.377124	0.038998	
free sulfur dioxide	-0.282735	-0.352557	0.133126	
total sulfur dioxide	-0.329054	-0.414476	0.195242	
density	0.458910	0.271296	0.096154	
pH	-0.252700	0.261454	-0.329808	
sulphates	0.299568	0.225984	0.056197	
alcohol	-0.095452	-0.037640	-0.010493	

	residual sugar	chlorides	free sulfur
dioxide \			
fixed acidity	-0.111981	0.298195	-0.282735
volatile acidity	-0.196011	0.377124	-0.352557
citric acid	0.142451	0.038998	0.133126

residual sugar	1.000000	-0.128940	0.402871
chlorides	-0.128940	1.000000	-0.195045
free sulfur dioxide	0.402871	-0.195045	1.000000
total sulfur dioxide	0.495482	-0.279630	0.720934
density	0.552517	0.362615	0.025717
pH	-0.267320	0.044708	-0.145854
sulphates	-0.185927	0.395593	-0.188457
alcohol	-0.359415	-0.256916	-0.179838

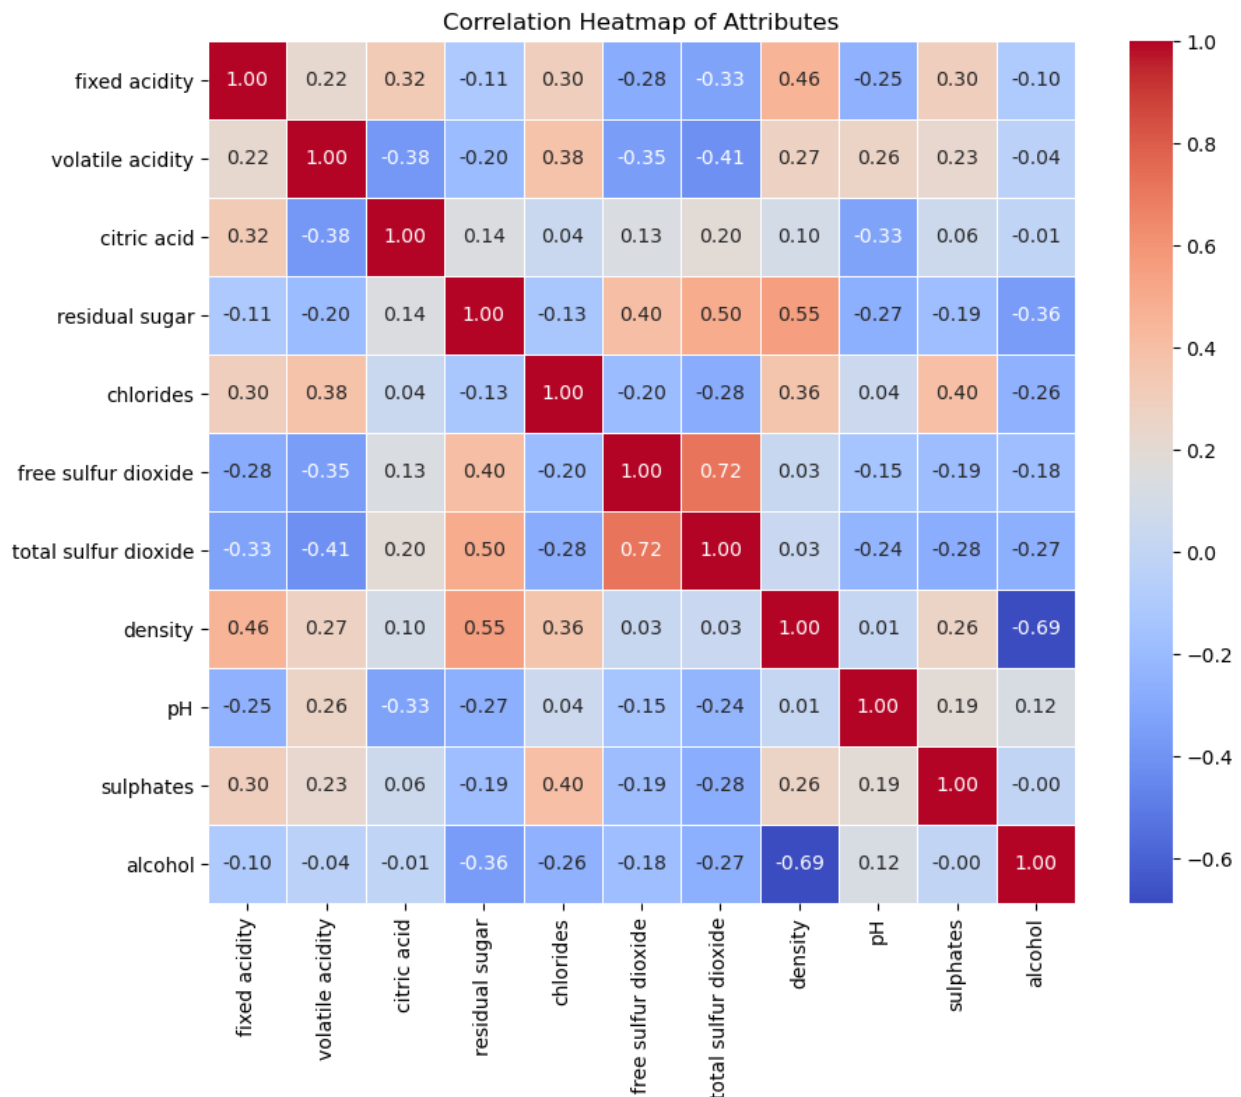
	total sulfur dioxide	density	pH	
sulphates \				
fixed acidity	-0.329054	0.458910	-0.252700	
0.299568				
volatile acidity	-0.414476	0.271296	0.261454	
0.225984				
citric acid	0.195242	0.096154	-0.329808	
0.056197				
residual sugar	0.495482	0.552517	-0.267320	-
0.185927				
chlorides	-0.279630	0.362615	0.044708	
0.395593				
free sulfur dioxide	0.720934	0.025717	-0.145854	-
0.188457				
total sulfur dioxide	1.000000	0.032395	-0.238413	-
0.275727				
density	0.032395	1.000000	0.011686	
0.259478				
pH	-0.238413	0.011686	1.000000	
0.192123				
sulphates	-0.275727	0.259478	0.192123	
1.000000				
alcohol	-0.265740	-0.686745	0.121248	-
0.003029				

	alcohol
fixed acidity	-0.095452
volatile acidity	-0.037640
citric acid	-0.010493
residual sugar	-0.359415
chlorides	-0.256916
free sulfur dioxide	-0.179838

```
total sulfur dioxide -0.265740
density             -0.686745
pH                  0.121248
sulphates           -0.003029
alcohol             1.000000
```

```
# Create a heatmap
```

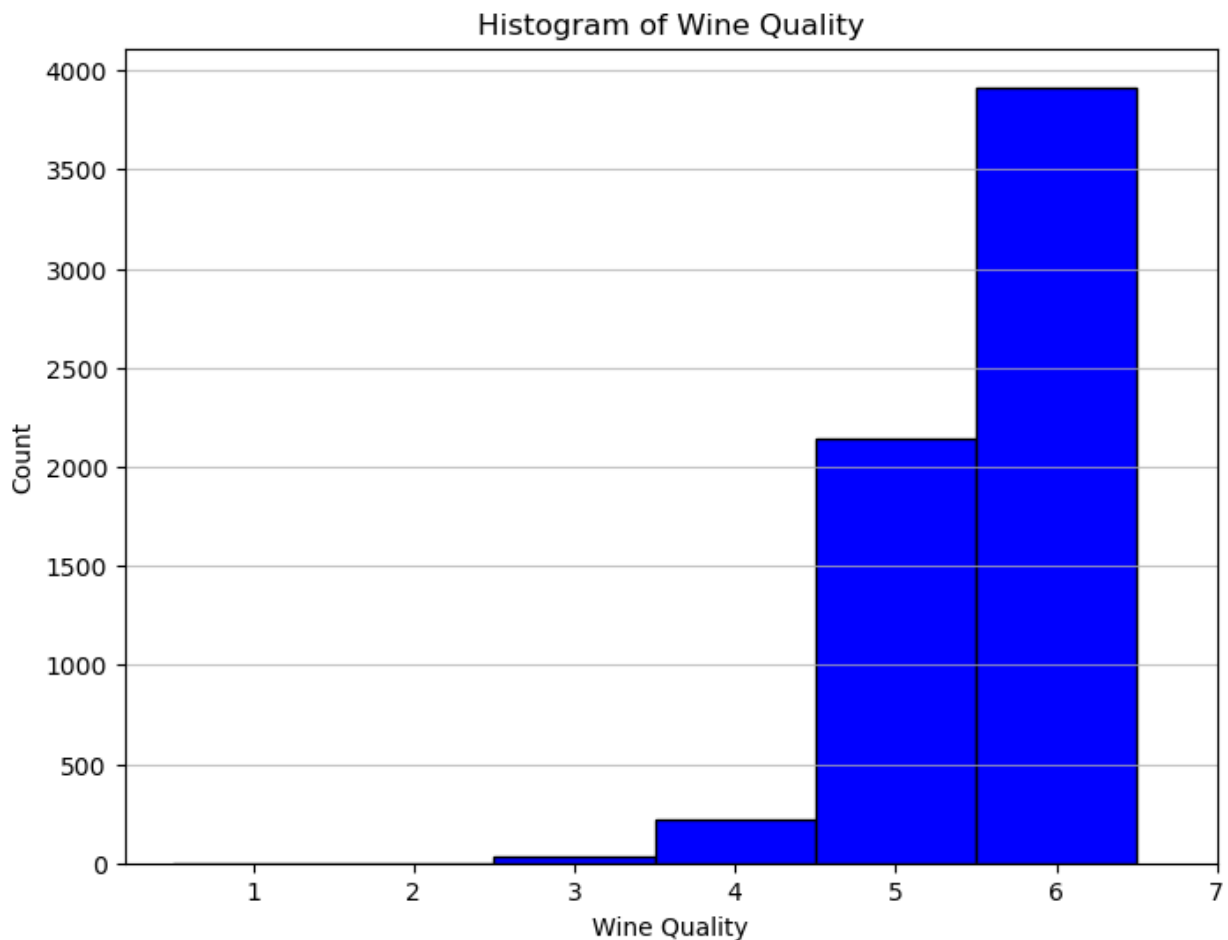
```
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm',
            fmt=".2f", linewidths=.5)
plt.title('Correlation Heatmap of Attributes')
plt.show()
```



## Question 5

Plot the frequency distribution of wine quality by using the Quality attribute. (8 points)

```
# Plot the histogram of wine quality
plt.figure(figsize=(8, 6))
plt.hist(wine_data_all['quality'], bins=range(1, 8),
         edgecolor='black', align='left', color='blue')
plt.title('Histogram of Wine Quality')
plt.xlabel('Wine Quality')
plt.ylabel('Count')
plt.xticks(range(1, 8))
plt.grid(axis='y', alpha=0.75)
plt.show()
```



## Question 6

Reduce the levels of rating for quality to three levels, i.e., high(2), medium(1), and low(0). Assign the levels 3 and 4 to level 0; 5 and 6 to level 1; and 7, 8, and 9 to level 2. You can use either "high, medium, low" or equivalent numbers. (10 points)

```
# Map the quality levels to the new reduced levels
quality_mapping = {0: [3, 4], 1: [5, 6], 2: [7, 8, 9]}
wine_data_all['quality'] = wine_data_all['quality'].apply(
```



```

lambda x: next((k for k, v in
quality_mapping.items() if x in v), None))
wine_data_all.head()

```

	fixed acidity chlorides \	volatile acidity	citric acid	residual sugar
0	7.4	0.70	0.00	1.9
1	7.8	0.88	0.00	2.6
2	7.8	0.76	0.04	2.3
3	11.2	0.28	0.56	1.9
4	7.4	0.70	0.00	1.9

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates
0	11.0	34.0	0.9978	3.51	0.56
1	25.0	67.0	0.9968	3.20	0.68
2	15.0	54.0	0.9970	3.26	0.65
3	17.0	60.0	0.9980	3.16	0.58
4	11.0	34.0	0.9978	3.51	0.56

	alcohol	quality
0	9.4	1
1	9.8	1
2	9.8	1
3	9.8	1
4	9.4	1

## Question 7

Normalize the numeric attributes. Hint:  $(x - \min(x)) / (\max(x) - \min(x))$  (10 points)

```

from sklearn.preprocessing import MinMaxScaler

numeric_attributes = wine_data_all.drop(['quality'], axis=1)

# Initialize the MinMaxScaler
scaler = MinMaxScaler()

# Fit and transform the numeric attributes using MinMaxScaler

```

```
normalized_data = scaler.fit_transform(numeric_attributes)

# Create a DataFrame with normalized values
normalized_df = pd.DataFrame(normalized_data,
                              columns=numeric_attributes.columns)

normalized_df.head()
```

	fixed acidity	volatile acidity	citric acid	residual sugar
0	0.297521	0.413333	0.000000	0.019939
1	0.330579	0.533333	0.000000	0.030675
2	0.330579	0.453333	0.024096	0.026074
3	0.611570	0.133333	0.337349	0.019939
4	0.297521	0.413333	0.000000	0.019939

	free sulfur dioxide	total sulfur dioxide	density	pH
0	0.034722	0.064516	0.206092	0.612403
1	0.083333	0.140553	0.186813	0.372093
2	0.048611	0.110599	0.190669	0.418605
3	0.055556	0.124424	0.209948	0.341085
4	0.034722	0.064516	0.206092	0.612403

	alcohol
0	0.202899
1	0.260870
2	0.260870
3	0.260870
4	0.202899

## Question 8

Divide the dataset to training and test sets. (10 points)

```
from sklearn.model_selection import train_test_split

# Separate features X and target variable y
X = wine_data_all.drop('quality', axis=1)
y = wine_data_all['quality']
```

```

# Split the dataset into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Display the shapes of the training and test sets
print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)

X_train shape: (5197, 11)
X_test shape: (1300, 11)
y_train shape: (5197,)
y_test shape: (1300,)

```

## Question 9

Use the Logistic Regression algorithm to predict the quality of wine using its attributes. (12 points)

```

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix

# Initialize the Logistic Regression model
model = LogisticRegression(max_iter=10000, solver='sag')

# Fit the model on the training data
model.fit(X_train, y_train)

# Make predictions on the test data
y_pred = model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred, zero_division=1)

# Display the results
print("Accuracy:", accuracy)
print("\nConfusion Matrix:", conf_matrix)
print("\nClassification Report:", class_report)

Accuracy: 0.7823076923076923

Confusion Matrix: [[ 0  49  0]
 [ 0 966 33]
 [ 0 201 51]]

```

Classification Report:			precision	recall	f1-score
support					
0	1.00	0.00	0.00	49	
1	0.79	0.97	0.87	999	
2	0.61	0.20	0.30	252	
accuracy			0.78	1300	
macro avg			0.39	0.39	1300
weighted avg			0.77	0.73	1300

## Question 10

Use the KNN algorithm to predict the quality of wine using its attributes. (12 points)

```
from sklearn.neighbors import KNeighborsClassifier

# Initialize the KNN model
model_knn = KNeighborsClassifier(n_neighbors=5)

# Fit the model on the training data
model_knn.fit(X_train, y_train)

# Make predictions on the test data
y_pred_knn = model_knn.predict(X_test)

# Evaluate the model
accuracy_knn = accuracy_score(y_test, y_pred_knn)
conf_matrix_knn = confusion_matrix(y_test, y_pred_knn)
class_report_knn = classification_report(y_test, y_pred_knn )

# Display the results
print("Accuracy:", accuracy_knn)
print("\nConfusion Matrix:", conf_matrix_knn)
print("\nClassification Report:", class_report_knn)
```

Accuracy: 0.7530769230769231

Confusion Matrix: [[ 3 44 2]  
[ 11 902 86]  
[ 1 177 74]]

Classification Report:			precision	recall	f1-score
support					
0	0.20	0.06	0.09	49	
1	0.80	0.90	0.85	999	
2	0.46	0.29	0.36	252	

accuracy			0.75	1300
macro avg	0.49	0.42	0.43	1300
weighted avg	0.71	0.75	0.73	1300

## Question 11

Display two confusion matrices to evaluate the performances of Logistic Regression and KNN. (A simple matrix is enough. No need to plot it.) (12 points)

```
from sklearn.metrics import confusion_matrix

# confusion matrix using Logistic Regression model
conf_matrix = confusion_matrix(y_test, y_pred)

# confusion matrix using KNN model
conf_matrix_knn = confusion_matrix(y_test, y_pred_knn)

print("Confusion Matrix for Logistic Regression:")
print(conf_matrix)

print("\nConfusion Matrix for KNN:")
print(conf_matrix_knn)
```

Confusion Matrix for Logistic Regression:

```
[[ 0  49   0]
 [ 0 966  33]
 [ 0 201  51]]
```

Confusion Matrix for KNN:

```
[[ 3  44   2]
 [11 902  86]
 [ 1 177  74]]
```

## Question 12

Evaluate the models' performances by computing Accuracy, Precision, and Recall. If you are using a package to calculate the values, you should explain what you understand from the output of the package by mentioning the exact accuracy, precision, and recall values in your own words. (12 points)

```
from sklearn.metrics import accuracy_score, precision_score,
recall_score

# Logistic Regression
logreg_accuracy = accuracy_score(y_test, y_pred)
logreg_precision = precision_score(y_test, y_pred, average='weighted',
                                   zero_division=1)
logreg_recall = recall_score(y_test, y_pred, average='weighted')
```

```

print("Logistic Regression Performance:")
print(f"Accuracy: {logreg_accuracy:.4f}")
print(f"Precision: {logreg_precision:.4f}")
print(f"Recall: {logreg_recall:.4f}\n")

# KNN
knn_accuracy = accuracy_score(y_test, y_pred_knn)
knn_precision = precision_score(y_test, y_pred_knn,
                                average='weighted')
knn_recall = recall_score(y_test, y_pred_knn, average='weighted')

print("KNN Performance:")
print(f"Accuracy: {knn_accuracy:.4f}")
print(f"Precision: {knn_precision:.4f}")
print(f"Recall: {knn_recall:.4f}")

Logistic Regression Performance:
Accuracy: 0.7823
Precision: 0.7659
Recall: 0.7823

KNN Performance:
Accuracy: 0.7531
Precision: 0.7133
Recall: 0.7531

```

- Accuracy: Accuracy tells us how many predictions a model got right. For Logistic Regression, it's about 78.23%, and for KNN, it's about 75.31%. So, Logistic Regression is a bit more accurate.
- Precision: Precision is about the accuracy of positive predictions. For Logistic Regression, when it says something is positive, it's right about 76.59% of the time. KNN is a bit lower, around 71.33%. Higher precision is better, meaning the model is more careful about saying things are positive.
- Recall: Recall is how well a model finds all the actual positive cases. For Logistic Regression, it captures about 78.23% of all actual positives, and for KNN, it's about 75.31%. Higher recall is better, especially in situations where missing positive cases is a big concern, like in medical diagnoses.
- Conclusion: Overall, Logistic Regression performs better in terms of accuracy, precision, and recall compared to KNN. So, for this dataset, Logistic Regression seems to be the better choice.

This is the end of Assignment 2

Ceni Babaoglu, PhD